

# hw4

May 21, 2021

```
[1]: import os
      #import cv2
      import numpy as np
      #from tqdm import tqdm
      import seaborn as sns
      import tensorflow as tf
      from tensorflow import keras
      import matplotlib.pyplot as plt
      from tensorflow.keras import layers
      from sklearn.decomposition import PCA
      from tensorflow.keras import models
      from mpl_toolkits.axes_grid1 import ImageGrid
      from sklearn.metrics import confusion_matrix ,plot_confusion_matrix
      from tensorflow.keras.preprocessing.image import ImageDataGenerator

      epocs= 50
```

## 0.1 Let's Get some data!

So we started by just getting the fashion mnist dataset. Tensorflow had this build it, so it was very convenient. Because the autoencoder takes in the training images as the label and the input the training labels were only needed for the final 2-D mapping.

```
[2]: (train_images, train_labels), (test_images, test_labels) = keras.datasets.
      ↪fashion_mnist.load_data()
      class_names = ['T-shirt/top', 'Trouser', 'Pullover', 'Dress', 'Coat',
                     'Sandal', 'Shirt', 'Sneaker', 'Bag', 'Ankle boot']

      #train_labels = np.asarray(train_labels).astype('float32').reshape((-1,1))
      #test_labels = np.asarray(test_labels).astype('float32').reshape((-1,1))
```

## 0.2 Make Noisy images now

- <https://keras.io/examples/vision/autoencoder/>

These functions I took directly from the Keras Site. It just displays the images and it also adds noise to the image. I did change the type of noise however, and I was sure to give it 0 mean gaussian noise.

```
[3]: def noise(array):
    """
    Adds random noise to each image in the supplied array.
    """

    mean = 0.0    # some constant
    std = 1.0     # some constant (standard deviation)
    noisy_array = array + np.random.normal(mean, std, array.shape)

    return np.clip(noisy_array, 0.0, 1.0)

def display(array1, array2):
    """
    Displays ten random images from each one of the supplied arrays.
    """

    n = 10

    indices = np.random.randint(len(array1), size=n)
    images1 = array1[indices, :]
    images2 = array2[indices, :]

    plt.figure(figsize=(20, 4))
    for i, (image1, image2) in enumerate(zip(images1, images2)):
        ax = plt.subplot(2, n, i + 1)
        plt.imshow(image1.reshape(28, 28))
        plt.gray()
        ax.get_xaxis().set_visible(False)
        ax.get_yaxis().set_visible(False)

        ax = plt.subplot(2, n, i + 1 + n)
        plt.imshow(image2.reshape(28, 28))
        plt.gray()
        ax.get_xaxis().set_visible(False)
        ax.get_yaxis().set_visible(False)

    plt.show()

[4]: train_images = train_images.astype('float32')/255.
    test_images = test_images.astype('float32')/255.

    noisy_train_data = noise(train_images)
    noisy_test_data = noise(test_images)
```

### 0.3 Build Model

The middle of the autoencoder was definitely the most difficult part. If I didn't get the sizes right, I ended up with an output image that was not  $(28 \times 28)$  and my final error check would fail. Though it's not 100% symmetric, it does line the final shape up very well.

The loss, stayed relatively close. It dropped initially and then stayed around  $\sim 0.39$ . With more training though my results would sometimes be better, other times, I would end up with just a white square.

```
[5]: input_img = keras.Input(shape=(28, 28, 1))
x = layers.Conv2D(16, (3, 3), activation='relu', padding='same')(input_img)
x = layers.MaxPooling2D((2, 2), padding='same')(x)
x = layers.Conv2D(8, (3, 3), activation='relu', padding='same')(x)
x = layers.MaxPooling2D((2, 2), padding='same')(x)
x = layers.Conv2D(8, (3, 3), activation='relu', padding='same')(x)
x = layers.MaxPooling2D((2, 2), padding='same')(x)
x = layers.Conv2D(4, (3, 3), activation='relu', padding='same')(x)
encoded = layers.MaxPooling2D((2, 2), padding='same')(x)

# at this point the representation is (4, 4, 8) i.e. 128-dimensional
bottleNeckOut = layers.Flatten(input_shape=x.shape)(encoded)
#bottleNeckOut = layers.Dense(4, activation='relu')(bottleNeckOut)
bottleNeckOut = layers.Dense(2, activation='relu')(bottleNeckOut)

bottleNeckIn = layers.Dense(4, activation='relu')(bottleNeckOut)
bottleNeckIn = layers.Reshape((2,2,1))(bottleNeckIn)

x = layers.Conv2D(4, (3, 3), activation='relu', padding='same')(bottleNeckIn)
x = layers.UpSampling2D((2, 2))(x)
x = layers.Conv2D(8, (3, 3), activation='relu', padding='same')(x)
x = layers.UpSampling2D((2, 2))(x)
x = layers.Conv2D(8, (3, 3), activation='relu', padding='same')(x)
x = layers.UpSampling2D((2, 2))(x)
x = layers.Conv2D(16, (3, 3), activation='relu')(x)
x = layers.UpSampling2D((2, 2))(x)
decoded = layers.Conv2D(1, (3, 3), activation='sigmoid', padding='same')(x)

autoencoder = keras.Model(input_img, decoded)
denoiser = keras.models.clone_model(autoencoder)
encoder = keras.Model(input_img, bottleNeckOut)
autoencoder.compile(optimizer='adam',
    ↪loss='binary_crossentropy',metrics=['accuracy'])
denoiser.compile(optimizer='adam',
    ↪loss='binary_crossentropy',metrics=['accuracy'])

autoencoder.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 28, 28, 1)]	0
conv2d (Conv2D)	(None, 28, 28, 16)	160
max_pooling2d (MaxPooling2D)	(None, 14, 14, 16)	0
conv2d_1 (Conv2D)	(None, 14, 14, 8)	1160
max_pooling2d_1 (MaxPooling2D)	(None, 7, 7, 8)	0
conv2d_2 (Conv2D)	(None, 7, 7, 8)	584
max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 8)	0
conv2d_3 (Conv2D)	(None, 4, 4, 4)	292
max_pooling2d_3 (MaxPooling2D)	(None, 2, 2, 4)	0
flatten (Flatten)	(None, 16)	0
dense (Dense)	(None, 2)	34
dense_1 (Dense)	(None, 4)	12
reshape (Reshape)	(None, 2, 2, 1)	0
conv2d_4 (Conv2D)	(None, 2, 2, 4)	40
up_sampling2d (UpSampling2D)	(None, 4, 4, 4)	0
conv2d_5 (Conv2D)	(None, 4, 4, 8)	296
up_sampling2d_1 (UpSampling2D)	(None, 8, 8, 8)	0
conv2d_6 (Conv2D)	(None, 8, 8, 8)	584
up_sampling2d_2 (UpSampling2D)	(None, 16, 16, 8)	0
conv2d_7 (Conv2D)	(None, 14, 14, 16)	1168
up_sampling2d_3 (UpSampling2D)	(None, 28, 28, 16)	0
conv2d_8 (Conv2D)	(None, 28, 28, 1)	145

Total params: 4,475  
Trainable params: 4,475  
Non-trainable params: 0

```
[6]: history = autoencoder.fit(train_images, train_images,  
                               batch_size=64,  
                               steps_per_epoch = 500,  
                               shuffle = True,  
                               epochs=epochs)
```

```
Epoch 1/50  
500/500 [=====] - 11s 21ms/step - loss: 0.5247 -  
accuracy: 0.4878  
Epoch 2/50  
500/500 [=====] - 11s 21ms/step - loss: 0.3830 -  
accuracy: 0.4882  
Epoch 3/50  
500/500 [=====] - 11s 21ms/step - loss: 0.3682 -  
accuracy: 0.4900  
Epoch 4/50  
500/500 [=====] - 11s 21ms/step - loss: 0.3577 -  
accuracy: 0.4938  
Epoch 5/50  
500/500 [=====] - 11s 21ms/step - loss: 0.3506 -  
accuracy: 0.4954  
Epoch 6/50  
500/500 [=====] - 11s 21ms/step - loss: 0.3497 -  
accuracy: 0.4941  
Epoch 7/50  
500/500 [=====] - 11s 22ms/step - loss: 0.3486 -  
accuracy: 0.4927  
Epoch 8/50  
500/500 [=====] - 11s 22ms/step - loss: 0.3447 -  
accuracy: 0.4970  
Epoch 9/50  
500/500 [=====] - 11s 21ms/step - loss: 0.3438 -  
accuracy: 0.4971  
Epoch 10/50  
500/500 [=====] - 11s 22ms/step - loss: 0.3442 -  
accuracy: 0.4948  
Epoch 11/50  
500/500 [=====] - 11s 22ms/step - loss: 0.3429 -  
accuracy: 0.4952  
Epoch 12/50  
500/500 [=====] - 11s 22ms/step - loss: 0.3434 -  
accuracy: 0.4947  
Epoch 13/50
```

500/500 [=====] - 11s 22ms/step - loss: 0.3396 -  
 accuracy: 0.4982  
 Epoch 14/50  
 500/500 [=====] - 11s 22ms/step - loss: 0.3404 -  
 accuracy: 0.4973  
 Epoch 15/50  
 500/500 [=====] - 11s 22ms/step - loss: 0.3415 -  
 accuracy: 0.4949  
 Epoch 16/50  
 500/500 [=====] - 11s 22ms/step - loss: 0.3395 -  
 accuracy: 0.4970  
 Epoch 17/50  
 500/500 [=====] - 11s 22ms/step - loss: 0.3398 -  
 accuracy: 0.4962  
 Epoch 18/50  
 500/500 [=====] - 11s 22ms/step - loss: 0.3385 -  
 accuracy: 0.4976  
 Epoch 19/50  
 500/500 [=====] - 11s 22ms/step - loss: 0.3381 -  
 accuracy: 0.4976  
 Epoch 20/50  
 500/500 [=====] - 11s 22ms/step - loss: 0.3380 -  
 accuracy: 0.4973  
 Epoch 21/50  
 500/500 [=====] - 11s 22ms/step - loss: 0.3380 -  
 accuracy: 0.4972  
 Epoch 22/50  
 500/500 [=====] - 11s 22ms/step - loss: 0.3372 -  
 accuracy: 0.4969  
 Epoch 23/50  
 500/500 [=====] - 11s 22ms/step - loss: 0.3369 -  
 accuracy: 0.4981  
 Epoch 24/50  
 500/500 [=====] - 11s 22ms/step - loss: 0.3359 -  
 accuracy: 0.4984  
 Epoch 25/50  
 500/500 [=====] - 11s 22ms/step - loss: 0.3376 -  
 accuracy: 0.4957  
 Epoch 26/50  
 500/500 [=====] - 11s 22ms/step - loss: 0.3374 -  
 accuracy: 0.4972  
 Epoch 27/50  
 500/500 [=====] - 11s 22ms/step - loss: 0.3368 -  
 accuracy: 0.4967  
 Epoch 28/50  
 500/500 [=====] - 11s 22ms/step - loss: 0.3367 -  
 accuracy: 0.4966  
 Epoch 29/50

500/500 [=====] - 11s 22ms/step - loss: 0.3371 -  
 accuracy: 0.4960  
 Epoch 30/50  
 500/500 [=====] - 11s 22ms/step - loss: 0.3358 -  
 accuracy: 0.4975  
 Epoch 31/50  
 500/500 [=====] - 11s 22ms/step - loss: 0.3359 -  
 accuracy: 0.4982  
 Epoch 32/50  
 500/500 [=====] - 11s 22ms/step - loss: 0.3366 -  
 accuracy: 0.4962  
 Epoch 33/50  
 500/500 [=====] - 11s 22ms/step - loss: 0.3364 -  
 accuracy: 0.4968  
 Epoch 34/50  
 500/500 [=====] - 11s 22ms/step - loss: 0.3349 -  
 accuracy: 0.4982  
 Epoch 35/50  
 500/500 [=====] - 11s 22ms/step - loss: 0.3352 -  
 accuracy: 0.4978  
 Epoch 36/50  
 500/500 [=====] - 11s 22ms/step - loss: 0.3342 -  
 accuracy: 0.4992  
 Epoch 37/50  
 500/500 [=====] - 11s 22ms/step - loss: 0.3355 -  
 accuracy: 0.4967  
 Epoch 38/50  
 500/500 [=====] - 11s 22ms/step - loss: 0.3351 -  
 accuracy: 0.4984  
 Epoch 39/50  
 500/500 [=====] - 11s 22ms/step - loss: 0.3341 -  
 accuracy: 0.4984  
 Epoch 40/50  
 500/500 [=====] - 11s 22ms/step - loss: 0.3353 -  
 accuracy: 0.4974  
 Epoch 41/50  
 500/500 [=====] - 11s 22ms/step - loss: 0.3343 -  
 accuracy: 0.4989  
 Epoch 42/50  
 500/500 [=====] - 11s 22ms/step - loss: 0.3332 -  
 accuracy: 0.4992  
 Epoch 43/50  
 500/500 [=====] - 11s 22ms/step - loss: 0.3347 -  
 accuracy: 0.4973  
 Epoch 44/50  
 500/500 [=====] - 11s 22ms/step - loss: 0.3345 -  
 accuracy: 0.4976  
 Epoch 45/50

```

500/500 [=====] - 11s 22ms/step - loss: 0.3338 -
accuracy: 0.4978
Epoch 46/50
500/500 [=====] - 11s 22ms/step - loss: 0.3338 -
accuracy: 0.4978
Epoch 47/50
500/500 [=====] - 11s 22ms/step - loss: 0.3338 -
accuracy: 0.4979
Epoch 48/50
500/500 [=====] - 11s 22ms/step - loss: 0.3339 -
accuracy: 0.4986
Epoch 49/50
500/500 [=====] - 11s 22ms/step - loss: 0.3326 -
accuracy: 0.4993
Epoch 50/50
500/500 [=====] - 11s 22ms/step - loss: 0.3339 -
accuracy: 0.4971

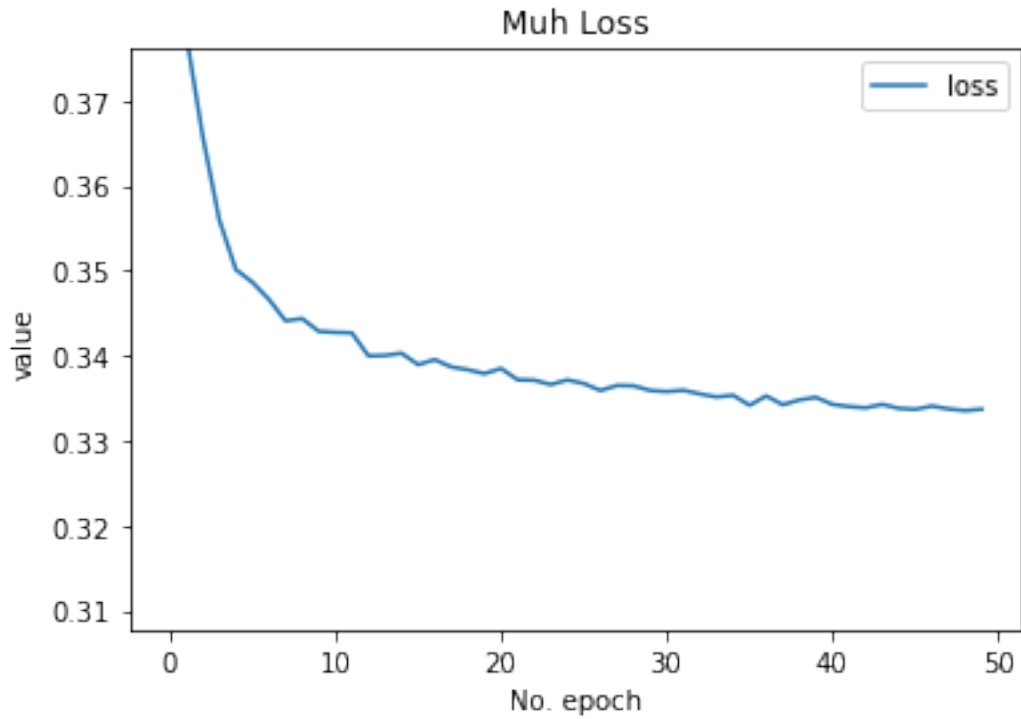
```

```

[7]: rebuild = autoencoder.predict(test_images)
avee = sum(history.history['loss'])/len(history.history['loss'])
plt.plot(history.history['loss'],label = "loss")
plt.title('Muh Loss')
plt.ylabel('value')
plt.xlabel('No. epoch')
plt.ylim(avee-(avee*.10),avee+(avee*.10))
plt.legend(loc="upper right")
plt.show()


```





```
[8]: n = 10
plt.figure(figsize=(20, 4))
for i in range(n):
    # display original
    ax = plt.subplot(2, n, i + 1)
    plt.imshow(test_images[i])
    plt.title("original")
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)

    # display reconstruction
    ax = plt.subplot(2, n, i + 1 + n)
    plt.imshow(rebuild[i])
    plt.title("reconstructed")
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)
plt.show()
```




## 1 De-noise

- <https://keras.io/examples/vision/autoencoder/>

My denoiser thinks everything looks like a shirt. I assume this is because shirts are rectangular and take up most of the image size. This is similar to what the noise looks like on the picture. If I use a different kind of noise, or less than a standard deviation of 1, my images come out a lot better. But I was trying to follow the instructions as stated.

```
[9]: display(train_images, noisy_train_data)
```



```
[10]: history = denoiser.fit(noisy_train_data, train_images,
                             batch_size=64,
                             steps_per_epoch = 500,
                             shuffle = True,
                             epochs=epochs)
```

Epoch 1/50

500/500 [=====] - 11s 22ms/step - loss: 0.5842 - accuracy: 0.4910

Epoch 2/50

500/500 [=====] - 11s 22ms/step - loss: 0.4960 - accuracy: 0.4834

Epoch 3/50

500/500 [=====] - 11s 22ms/step - loss: 0.4942 - accuracy: 0.4853

Epoch 4/50

500/500 [=====] - 11s 22ms/step - loss: 0.4930 -

accuracy: 0.4858  
Epoch 5/50  
500/500 [=====] - 11s 22ms/step - loss: 0.4918 -  
accuracy: 0.4875  
Epoch 6/50  
500/500 [=====] - 11s 22ms/step - loss: 0.4920 -  
accuracy: 0.4866  
Epoch 7/50  
500/500 [=====] - 11s 22ms/step - loss: 0.4921 -  
accuracy: 0.4864  
Epoch 8/50  
500/500 [=====] - 11s 22ms/step - loss: 0.4910 -  
accuracy: 0.4877  
Epoch 9/50  
500/500 [=====] - 11s 22ms/step - loss: 0.4922 -  
accuracy: 0.4875  
Epoch 10/50  
500/500 [=====] - 11s 22ms/step - loss: 0.4911 -  
accuracy: 0.4873  
Epoch 11/50  
500/500 [=====] - 11s 22ms/step - loss: 0.4911 -  
accuracy: 0.4898  
Epoch 12/50  
500/500 [=====] - 11s 22ms/step - loss: 0.4909 -  
accuracy: 0.4878  
Epoch 13/50  
500/500 [=====] - 11s 22ms/step - loss: 0.4910 -  
accuracy: 0.4881  
Epoch 14/50  
500/500 [=====] - 11s 22ms/step - loss: 0.4920 -  
accuracy: 0.4857  
Epoch 15/50  
500/500 [=====] - 11s 22ms/step - loss: 0.4915 -  
accuracy: 0.4867  
Epoch 16/50  
500/500 [=====] - 11s 22ms/step - loss: 0.4913 -  
accuracy: 0.4879  
Epoch 17/50  
500/500 [=====] - 11s 22ms/step - loss: 0.4903 -  
accuracy: 0.4876  
Epoch 18/50  
500/500 [=====] - 11s 22ms/step - loss: 0.4912 -  
accuracy: 0.4877  
Epoch 19/50  
500/500 [=====] - 11s 22ms/step - loss: 0.4917 -  
accuracy: 0.4879  
Epoch 20/50  
500/500 [=====] - 11s 22ms/step - loss: 0.4911 -

accuracy: 0.4872  
Epoch 21/50  
500/500 [=====] - 11s 22ms/step - loss: 0.4901 -  
accuracy: 0.4887  
Epoch 22/50  
500/500 [=====] - 11s 22ms/step - loss: 0.4917 -  
accuracy: 0.4878  
Epoch 23/50  
500/500 [=====] - 11s 22ms/step - loss: 0.4913 -  
accuracy: 0.4873  
Epoch 24/50  
500/500 [=====] - 11s 22ms/step - loss: 0.4903 -  
accuracy: 0.4882  
Epoch 25/50  
500/500 [=====] - 11s 22ms/step - loss: 0.4913 -  
accuracy: 0.4865  
Epoch 26/50  
500/500 [=====] - 11s 22ms/step - loss: 0.4916 -  
accuracy: 0.4865  
Epoch 27/50  
500/500 [=====] - 11s 22ms/step - loss: 0.4899 -  
accuracy: 0.4893  
Epoch 28/50  
500/500 [=====] - 11s 22ms/step - loss: 0.4911 -  
accuracy: 0.4882  
Epoch 29/50  
500/500 [=====] - 11s 22ms/step - loss: 0.4915 -  
accuracy: 0.4867  
Epoch 30/50  
500/500 [=====] - 11s 22ms/step - loss: 0.4915 -  
accuracy: 0.4872  
Epoch 31/50  
500/500 [=====] - 11s 22ms/step - loss: 0.4913 -  
accuracy: 0.4864  
Epoch 32/50  
500/500 [=====] - 11s 22ms/step - loss: 0.4898 -  
accuracy: 0.4887  
Epoch 33/50  
500/500 [=====] - 11s 22ms/step - loss: 0.4901 -  
accuracy: 0.4882  
Epoch 34/50  
500/500 [=====] - 11s 22ms/step - loss: 0.4909 -  
accuracy: 0.4878  
Epoch 35/50  
500/500 [=====] - 11s 22ms/step - loss: 0.4912 -  
accuracy: 0.4864  
Epoch 36/50  
500/500 [=====] - 11s 22ms/step - loss: 0.4918 -

```


accuracy: 0.4867
Epoch 37/50
500/500 [=====] - 11s 22ms/step - loss: 0.4908 -
accuracy: 0.4873
Epoch 38/50
500/500 [=====] - 11s 22ms/step - loss: 0.4913 -
accuracy: 0.4863
Epoch 39/50
500/500 [=====] - 11s 22ms/step - loss: 0.4903 -
accuracy: 0.4881
Epoch 40/50
500/500 [=====] - 11s 22ms/step - loss: 0.4917 -
accuracy: 0.4869
Epoch 41/50
500/500 [=====] - 11s 22ms/step - loss: 0.4915 -
accuracy: 0.4866
Epoch 42/50
500/500 [=====] - 11s 22ms/step - loss: 0.4914 -
accuracy: 0.4871
Epoch 43/50
500/500 [=====] - 11s 22ms/step - loss: 0.4914 -
accuracy: 0.4870
Epoch 44/50
500/500 [=====] - 11s 22ms/step - loss: 0.4904 -
accuracy: 0.4879
Epoch 45/50
500/500 [=====] - 11s 22ms/step - loss: 0.4906 -
accuracy: 0.4887
Epoch 46/50
500/500 [=====] - 11s 22ms/step - loss: 0.4914 -
accuracy: 0.4861
Epoch 47/50
500/500 [=====] - 11s 22ms/step - loss: 0.4905 -
accuracy: 0.4885
Epoch 48/50
500/500 [=====] - 11s 22ms/step - loss: 0.4910 -
accuracy: 0.4880
Epoch 49/50
500/500 [=====] - 11s 22ms/step - loss: 0.4897 -
accuracy: 0.4893
Epoch 50/50
500/500 [=====] - 11s 22ms/step - loss: 0.4905 -
accuracy: 0.4877

```

```

[11]: predictions = autoencoder.predict(noisy_test_data)
      display(noisy_test_data, predictions)

```




## 2 2-D Plot

The 2-D plot didn't always come out nice. There's definitely some randomness to the 2-D plot. I realize that sometimes I get a nice shape that is more similar to the PCA, but other times, it will be very far off.

```
[13]: cdict = {0: 'pink', 1: 'red', 2: 'blue', 3: 'green', 4: 'orange', 5: 'magenta', 6:
      ↪ 'olive',
      7: 'cyan', 8: 'purple', 9: 'yellow'}
color_list = [cdict[i] for i in test_labels ]
x_test_encoded = encoder.predict(test_images, batch_size=32)
print(x_test_encoded.shape)
plt.figure(figsize=(16, 10))
plt.scatter(x_test_encoded[:, 0], x_test_encoded[:, 1], c = color_list)
plt.show()
```

(10000, 2)



### 3 PCA


For the PCA portion. I didn't understand what Dr. Rhodes meant by "using a for loop", so I instead just did this in the same way that've used PCA in my last classes. I transformed the data from 3 dimensions to 2 before I fed it into the PCA function. In otherwords, I just vectorized the images using reshape. In this way I could reduce the components down to just two parameters(2 columns) and make the 2-D graph from that.

The shape of the PCA and the autoencoder is not always the same, but it is similar. There's definitely an overall shape that comes up when this is trained. I assume with more training the autoencoder and PCA 2-D maps would look more similar. But I was having an issue with consistency in my autoencoder.

```
[14]: pca = PCA(n_components=2)
      ti = test_images.reshape(10000,-1)
      principalComponents = pca.fit_transform(ti)
      principalComponents.shape
```

```
[14]: (10000, 2)
```

```
[15]: m,n = principalComponents.shape
      plt.figure(figsize=(16, 10))
      plt.scatter(principalComponents[:, 0], principalComponents[:, 1], c = ↵
      ↵color_list[:m])
      plt.show()
```



[ ]: